

# Predicting Airline Passenger Satisfaction

---

Airin Konno



# Aim

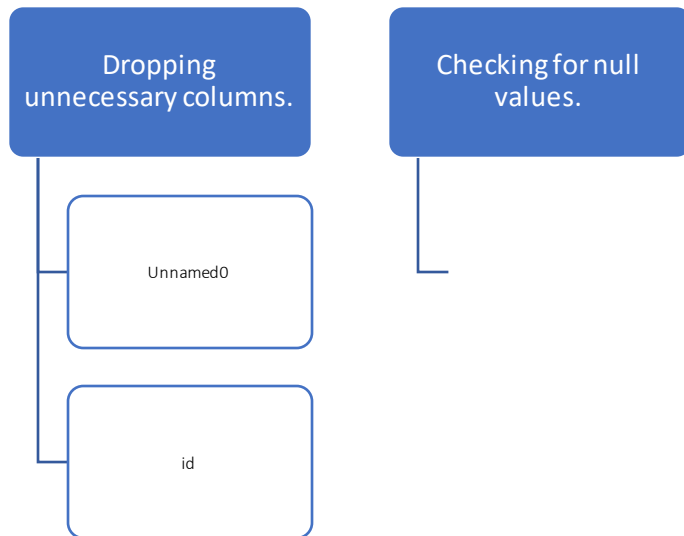
- Given survey data from an airline, predict customer satisfaction (satisfied or dissatisfied) using ML algorithms.
- Conduct EDA and data visualization to better understand the data.
- Potential business value:
  - Understand which features have the biggest impact on customer satisfaction and work towards improving those features.

# Data Collection





Dataset from Kaggle.

- Already split into train (80% of data) and test (20% of data) files.


# Data Preprocessing




# EDA – Categorical Variables Distribution



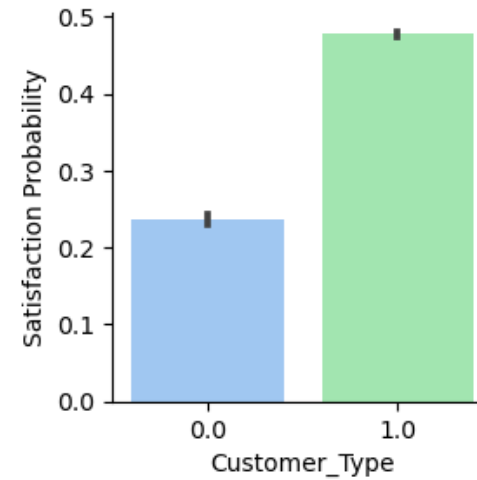
# EDA – Relationship with Satisfaction



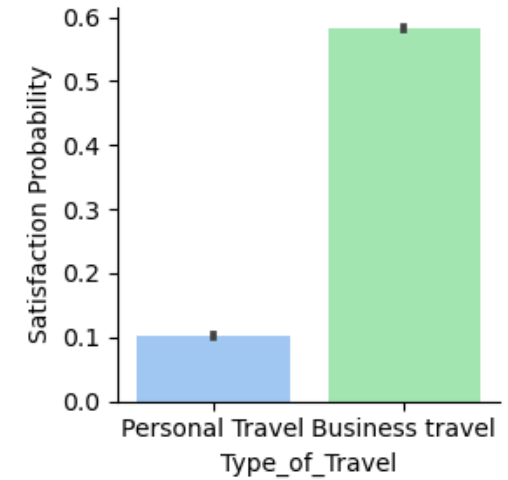
	Gender	satisfaction
1	Male	0.439475
0	Female	0.427371



	Class	satisfaction
2	2.0	0.694251
1	1.0	0.246064
0	0.0	0.186138


















	Customer_Type	satisfaction
1	1.0	0.477291
0	0.0	0.236658




	Type_of_Travel	satisfaction
0	Business travel	0.582597
1	Personal Travel	0.101678

# EDA – Numerical Variables Distribution



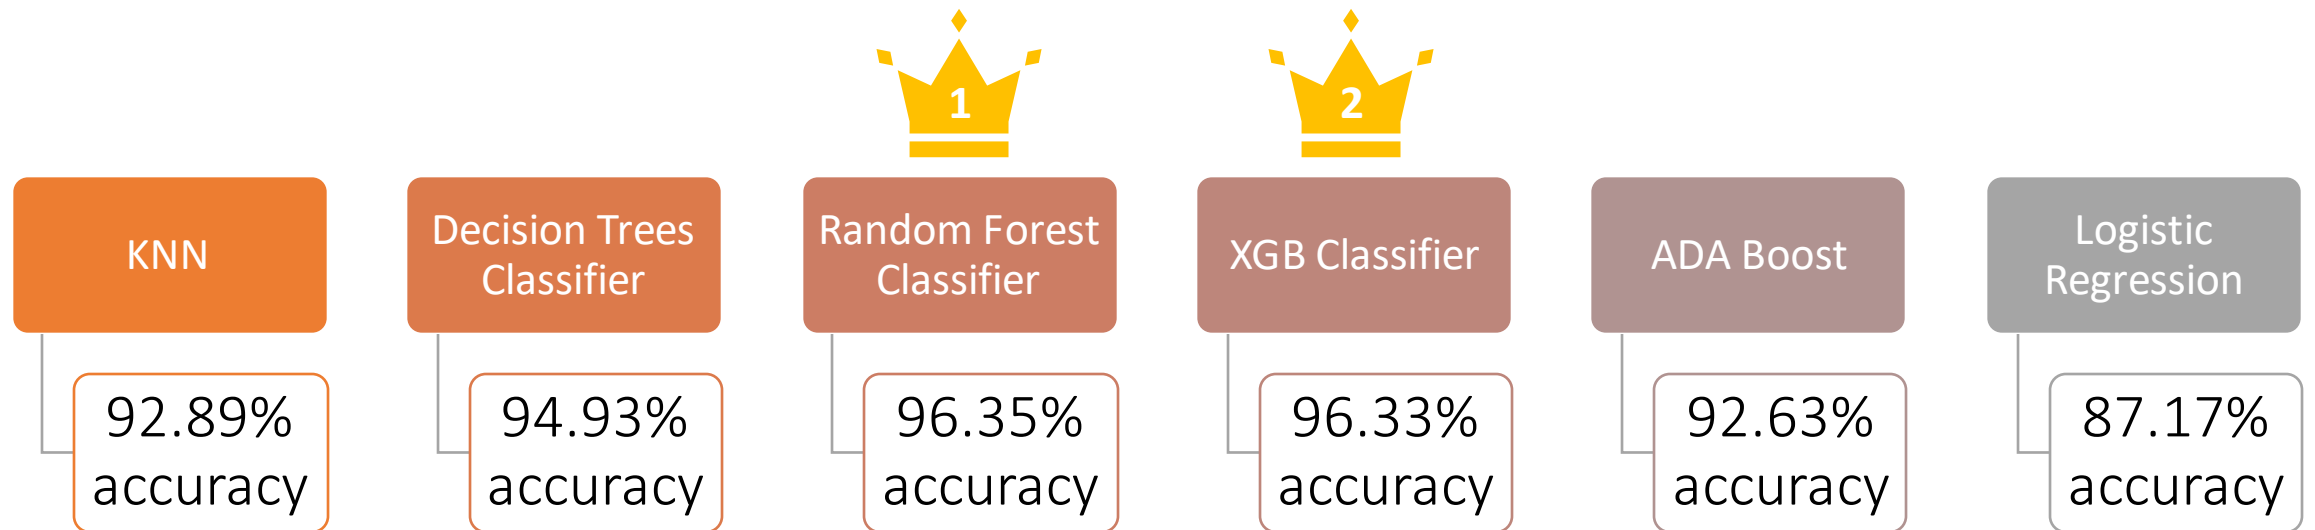
# EDA - Correlation Heat Map

- Shows which features correlate well with customer satisfaction.
- Best features
  - Online Booking, Class, and Inflight Entertainment
- Worst features
  - Type of Travel, Arrival Delay in Minutes, Departure/Arrival time convenient





# Machine Learning Algorithms





# Model 1



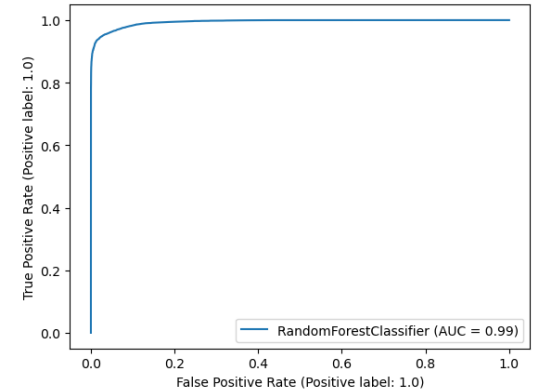

Random Forest  
Classifier



Accuracy:  
96.35%



Next steps: PCA,  
Grid Search



	precision	recall	f1-score	support
0.0	0.96	0.98	0.97	14573
1.0	0.97	0.94	0.96	11403
accuracy			0.96	25976
macro avg	0.96	0.96	0.96	25976
weighted avg	0.96	0.96	0.96	25976

# Model 1: Dimensionality Reduction PCA

## Grid Search

```
In [177]: param_grid = {  
    'n_estimators' : [int(x) for x in np.linspace(start = 800, stop = 1600, num = 3)],  
    'max_features' : ['auto', 'sqrt'],  
    'max_depth' : [int(x) for x in np.linspace(10, 110, num = 3)],  
    'min_samples_split' : [5, 10],  
    'min_samples_leaf' : [2, 4],  
    'bootstrap' : [True, False]  
}
```

```
In [178]: print(param_grid)  
  
{'n_estimators': [800, 1200, 1600], 'max_features': ['auto', 'sqrt'], 'max_depth': [10, 60, 110], 'min_samples_split': [5, 10], 'min_samples_leaf': [2, 4], 'bootstrap': [True, False]}
```

```
In [179]: # Create a based model  
rf_gd = RandomForestClassifier()
```

```
In [180]: grid_search = GridSearchCV(estimator = rf_gd, param_grid = param_grid,  
    cv = 2, n_jobs = -1, verbose = 2)
```

## Model 1: Grid Search

6 parameters were  
tuned.

There were a total  
of  
144 combinations.

Fitted 2 folds,  
creating 288 fits in  
total.

# Model 1:

## Grid Search Results

```
best_grid = grid_search.best_estimator_  
grid_accuracy = evaluate(best_grid, X_train, y_train)
```

Model Performance  
Accuracy is: 99.92300585155527 %

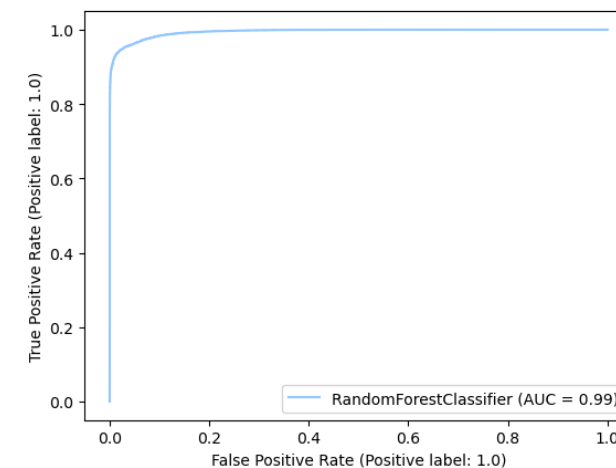

```
base_model = RandomForestClassifier(n_estimators = 1000)  
base_model.fit(X_train, y_train)  
base_accuracy = evaluate(base_model, X_test, y_test)
```

Model Performance  
Accuracy is: 96.34662765629812 %

```
print('Improvement of {:.2f}%'.format( 100 * (grid_accuracy - base_accuracy) / base_accuracy))  
Improvement of 3.71%.
```

best\_grid

```
RandomForestClassifier(bootstrap=False, max_depth=60, min_samples_leaf=2,  
                        min_samples_split=5, n_estimators=1600)
```





## Model 2





XGB Classifier



Accuracy: 96.33%



Next steps: PCA,  
Randomized Search



	precision	recall	f1-score	support
0.0	0.96	0.98	0.97	14573
1.0	0.97	0.94	0.96	11403
accuracy			0.96	25976
macro avg	0.96	0.96	0.96	25976
weighted avg	0.96	0.96	0.96	25976

# Model 2: Dimensionality Reduction PCA

## Randomized Search CV

```
In [200]: from sklearn.model_selection import RandomizedSearchCV
```

```
In [205]: classifier_xgb = xgb.XGBClassifier()
```

```
In [206]: params_rs_xgb = {  
    'learning_rate' : [0.05,0.10,0.15,0.20,0.25,0.30],  
    'max_depth' : [ 3, 4, 5, 6, 8, 10, 12, 15],  
    'min_child_weight' : [ 1, 3, 5, 7 ],  
    'gamma': [ 0.0, 0.1, 0.2 , 0.3, 0.4 ],  
    'colsample_bytree' : [ 0.3, 0.4, 0.5 , 0.7 ]  
}
```

```
In [207]: rs_model = RandomizedSearchCV(classifier_xgb,param_distributions=params_rs_xgb,n_iter=5,scoring='roc_auc')
```

```
In [208]: #model fitting  
rs_model.fit(X_train, y_train)
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

Model 1: Randomized  
Search

5 parameters  
were tuned.

Fitted 5 folds for  
each of the 5  
candidates

Total = 25 fits



# Model 2:

## Randomized Search Results

```
base_model_xgb = xgb.XGBClassifier()  
base_model_xgb.fit(X_train, y_train)  
base_accuracy_xgb = evaluate(base_model_xgb, X_test, y_test)
```

Model Performance  
Accuracy is: 96.33122882660918 %



```
best_random = rs_model.best_estimator_  
print(best_random)  
random_accuracy = evaluate(best_random, X_test, y_test)
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,  
              colsample_bylevel=None, colsample_bynode=None,  
              colsample_bytree=0.7, early_stopping_rounds=None,  
              enable_categorical=False, eval_metric=None, feature_types=None,  
              gamma=0.1, gpu_id=None, grow_policy=None, importance_type=None,  
              interaction_constraints=None, learning_rate=0.3, max_bin=None,  
              max_cat_threshold=None, max_cat_to_onehot=None,  
              max_delta_step=None, max_depth=8, max_leaves=None,  
              min_child_weight=5, missing=nan, monotone_constraints=None,  
              n_estimators=100, n_jobs=None, num_parallel_tree=None,  
              predictor=None, random_state=None, ...)
```

Model Performance  
Accuracy is: 96.29658145980905 %

```
print('Improvement of {:.2f}%'.format( 100 * (random_accuracy - base_accuracy_xgb) / base_accuracy_xgb))
```

Improvement of -0.04%.



# Modelling Conclusions

- Model 1 – Random Forest Classifier with Grid Search worked best (99.92% accuracy achieved).
- Even though PCA was conducted, I still proceeded with using the original X\_train and X\_test value.
- Randomized search also did not help Model 2.
  - Perhaps there were too few iterations.
  - Can try to increase the iterations next time to have more fits, which could help with the accuracy.